

a02-fnn_solution

May 5, 2024

1 Work done by pair of students

- Nico Sharei (nsharei)
- Artem Bisliouk (abisliou)

2 Deep Learning: Assignment 2

```
[ ]: # Define imports & defaults
import numpy as np
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset

# import helper functions
import os
import sys

sys.path.append(os.getcwd())
from a02helper import *

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
torch.manual_seed(0) # ensure reproducibility
np.random.seed(0)
MAX_SEQ_LEN = 200
BATCH_SIZE = 32
```

2.1 Task 1: Datasets

```
[ ]: import string
from torchtext.data import get_tokenizer

class ReviewsDataset(Dataset):
    def __init__(
        self,
        reviews_file="data/reviews_small.txt",
        labels_file="data/labels_small.txt",
        use_vocab=False,
```

```

):
    """
    A dataset of movie reviews and their labels.

    Args:
        reviews_file: the reviews file
        labels_file: the labels file
        use_vocab: if True, yield reviews in a numerical representation
    """
    # Load data from filesystem
    with open(reviews_file) as f:
        raw_reviews = f.readlines()

    with open(labels_file) as f:
        raw_labels = f.readlines()

    # Preprocessing and store (in memory)
    self._reviews = self._preprocess_reviews(raw_reviews)
    self._labels = self._preprocess_labels(raw_labels)

    # Build vocabulary
    self.vocab = None
    if use_vocab:
        from torchtext.vocab import build_vocab_from_iterator

        self.vocab = build_vocab_from_iterator(
            self._reviews, specials=["<pad>"] # will get token id 0
        )

    def __len__(self):
        """Returns the length of the dataset."""
        # YOUR CODE HERE
        return len(self._reviews)

    def __getitem__(self, idx):
        """
        Returns a tuple of a preprocessed review and the corresponding label.
        ↪ If the
        vocabulary is enabled, returns a numerical representation of the review.

        Args:
            idx: a single index
        Returns: a (review, label) tuple
        """
        # YOUR CODE HERE
        review, label = self._reviews[idx], self._labels[idx]

```

```

    # If vocabulary is enabled, convert review to numerical representation
    if self.vocab is not None:
        review = [self.vocab[token] for token in review]

    return review, label

def tokenize_review(self, review):
    tokenizer = get_tokenizer("basic_english")
    return tokenizer(review)

def _preprocess_reviews(self, raw_reviews):
    """
    Applies two kinds of preprocessing:

    (i) Apply the "basic_english" tokenizer from the torchtext library to
    transform every review into a list of normalized tokens (cf.
    https://pytorch.org/text/stable/data\_utils.html#get-tokenizer).

    (ii) Remove punctuation (cf.
    https://docs.python.org/3/library/string.html#string.punctuation).

    Returns: list of tokenized reviews
    """
    # YOUR CODE HERE

    tokenized_reviews = list(map(self.tokenize_review, raw_reviews))
    cleaned_reviews = [[token for token in review if token not in string.
↪punctuation] for review in tokenized_reviews]

    return cleaned_reviews

def _preprocess_labels(self, raw_labels):
    """
    Transform raw labels into integers, where 1="positive" and 0 otherwise.

    Returns: list of labels
    """
    # YOUR CODE HERE
    preprocessed_labels = [1 if label.strip() == 'positive' else 0 for ↪
↪label in raw_labels]

```

```

    return preprocessed_labels
    # Hint: You can remove leading and trailing whitespace from the raw
    ↪ labels using
    # the strip() method.

```

```
[ ]: # Test your code (without vocabulary).
```

```
dataset = ReviewsDataset()
```

```
print(dataset[0])
```

```
# Should yield:
```

```
# (['bromwell', 'high', 'is', 'a', 'cartoon', 'comedy', ... ], 1)
```

```

(['bromwell', 'high', 'is', 'a', 'cartoon', 'comedy', 'it', 'ran', 'at', 'the',
'same', 'time', 'as', 'some', 'other', 'programs', 'about', 'school', 'life',
'such', 'as', 'teachers', 'my', 'years', 'in', 'the', 'teaching', 'profession',
'lead', 'me', 'to', 'believe', 'that', 'bromwell', 'high', 's', 'satire', 'is',
'much', 'closer', 'to', 'reality', 'than', 'is', 'teachers', 'the', 'scramble',
'to', 'survive', 'financially', 'the', 'insightful', 'students', 'who', 'can',
'see', 'right', 'through', 'their', 'pathetic', 'teachers', 'pomp', 'the',
'pettiness', 'of', 'the', 'whole', 'situation', 'all', 'remind', 'me', 'of',
'the', 'schools', 'i', 'knew', 'and', 'their', 'students', 'when', 'i', 'saw',
'the', 'episode', 'in', 'which', 'a', 'student', 'repeatedly', 'tried', 'to',
'burn', 'down', 'the', 'school', 'i', 'immediately', 'recalled', 'at', 'high',
'a', 'classic', 'line', 'inspector', 'i', 'm', 'here', 'to', 'sack', 'one',
'of', 'your', 'teachers', 'student', 'welcome', 'to', 'bromwell', 'high', 'i',
'expect', 'that', 'many', 'adults', 'of', 'my', 'age', 'think', 'that',
'bromwell', 'high', 'is', 'far', 'fetched', 'what', 'a', 'pity', 'that', 'it',
'isn', 't'], 1)

```

```
[ ]: # Test your code (with vocabulary).
```

```
dataset = ReviewsDataset(use_vocab=True)
```

```
print(dataset[0])
```

```
# Should yield:
```

```
# ([10661, 307, 6, 3, 1177, 202, 8, ... ], 1)
```

```

([10661, 307, 6, 3, 1177, 202, 8, 2248, 33, 1, 168, 56, 15, 49, 85, 8902, 43,
422, 122, 140, 15, 3234, 59, 144, 9, 1, 5504, 6267, 454, 72, 5, 260, 12, 10661,
307, 13, 2060, 6, 73, 2780, 5, 692, 76, 6, 3234, 1, 29527, 5, 1730, 7117, 1,
6161, 1726, 36, 52, 68, 212, 143, 63, 1409, 3234, 17974, 1, 28056, 4, 1, 221,
758, 31, 2748, 72, 4, 1, 6311, 10, 731, 2, 63, 1726, 54, 10, 208, 1, 321, 9, 64,
3, 1601, 4042, 743, 5, 2853, 187, 1, 422, 10, 1254, 10116, 33, 307, 3, 380, 322,
6162, 10, 135, 136, 5, 10172, 30, 4, 134, 3234, 1601, 2545, 5, 10661, 307, 10,
529, 12, 113, 1841, 4, 59, 676, 103, 12, 10661, 307, 6, 227, 4163, 48, 3, 2201,
12, 8, 231, 21], 1)

```

2.2 Task 2: Data Loaders

```
[ ]: # Split dataset into training, validation, and test subsets
dataset = ReviewsDataset(use_vocab=True)
train_set, val_set, test_set = torch.utils.data.random_split(
    dataset, [0.8, 0.1, 0.1], generator=torch.Generator().manual_seed(123)
)
```

2.2.1 Task 2a

```
[ ]: # Example usage of a data loader
dataloader = DataLoader(
    val_set, # a dataset
    1, # desired batch size
    False, # whether to randomly shuffle the dataset
    num_workers=0, # number of workers that construct batches in parallel
)
```

```
[ ]: # Let's print the first batch
batch = next(iter(dataloader))
print(batch)

# [[tensor([11]), tensor([6]), tensor([1]), ..., tensor([0])]
```

```
[[tensor([11]), tensor([6]), tensor([1]), tensor([1037]), tensor([6578]),
tensor([4]), tensor([10]), tensor([89]), tensor([120]), tensor([48]),
tensor([163]), tensor([47]), tensor([6]), tensor([27]), tensor([342]),
tensor([4]), tensor([2228]), tensor([140]), tensor([3]), tensor([186]),
tensor([1466]), tensor([1]), tensor([771]), tensor([26]), tensor([78]),
tensor([1459]), tensor([200]), tensor([1101]), tensor([1]), tensor([66]),
tensor([26]), tensor([78]), tensor([5199]), tensor([5]), tensor([2288]),
tensor([1]), tensor([7861]), tensor([6591]), tensor([2]), tensor([83]),
tensor([2446]), tensor([25]), tensor([1]), tensor([182]), tensor([2756]),
tensor([2520]), tensor([34]), tensor([1]), tensor([145]), tensor([8]),
tensor([13]), tensor([39]), tensor([290]), tensor([25]), tensor([252]),
tensor([14480]), tensor([32]), tensor([52]), tensor([497]), tensor([9]),
tensor([223]), tensor([1]), tensor([3254]), tensor([25]), tensor([937]),
tensor([2]), tensor([153]), tensor([568]), tensor([5]), tensor([91]),
tensor([2]), tensor([30]), tensor([388]), tensor([1110]), tensor([17]),
tensor([80]), tensor([62]), tensor([1]), tensor([119]), tensor([255]),
tensor([14]), tensor([34]), tensor([2632]), tensor([1539]), tensor([133]),
tensor([28]), tensor([6]), tensor([31]), tensor([56]), tensor([33]),
tensor([27]), tensor([119]), tensor([413]), tensor([1]), tensor([230]),
tensor([4]), tensor([1036]), tensor([17]), tensor([3]), tensor([738]),
tensor([552]), tensor([1305]), tensor([11189]), tensor([14923]), tensor([6]),
tensor([88]), tensor([203]), tensor([3]), tensor([50]), tensor([283]),
tensor([19]), tensor([1]), tensor([27078]), tensor([44]), tensor([683]),
tensor([8]), tensor([0])]
```

2.2.2 Task 2b

```
[ ]: from torch.nn.utils.rnn import pad_sequence

def review_collate_fn(raw_batch):
    """Prepare batches of reviews from a review dataset.

    Args:
        raw_batch: collection of (review, label)-tuples from a ReviewDataset

    Returns: a tuple (review x token id tensor, label tensor) of sizes
        batch_size*MAX_SEQ_LEN and batch_size, respectively.

    """

    reviews = [item[0] for item in raw_batch]
    labels = [item[1] for item in raw_batch]

    cropped_reviews = [review[:MAX_SEQ_LEN] for review in reviews]

    padded_reviews = pad_sequence([torch.tensor(review) for review in
    ↵cropped_reviews], batch_first=True, padding_value=0)

    label_tensor = torch.tensor(labels)

    return padded_reviews, label_tensor
```

```
[ ]: # Test your function
review_collate_fn([([1, 2, 3], 1), (torch.arange(MAX_SEQ_LEN * 2) + 1, 0)])

# Should yield:
# (tensor([[ 1,  2,  3,  0,  0, ..., 0],
#          [ 1,  2,  3,  4,  5, ..., 200]]),
#  tensor([1, 0]))
```

```

/var/folders/t3/h38q5w_d36ncdxy42rj79mr0000gn/T/ipykernel_24909/2946447728.py:2
0: UserWarning: To copy construct from a tensor, it is recommended to use
sourceTensor.clone().detach() or
sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).

```

```

    padded_reviews = pad_sequence([torch.tensor(review) for review in
cropped_reviews], batch first=True, padding_value=0)

```

```
[ ]: (tensor([[ 1,  2,  3,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0]]))
```

```

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0],
[ 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14,
15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28,
29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42,
43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56,
57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70,
71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84,
85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98,
99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112,
113, 114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126,
127, 128, 129, 130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140,
141, 142, 143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154,
155, 156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168,
169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181, 182,
183, 184, 185, 186, 187, 188, 189, 190, 191, 192, 193, 194, 195, 196,
197, 198, 199, 200]]),
tensor([1, 0]))

```

```

[ ]: # Create the data loaders (with shuffling for training data -> randomization)
train_loader = DataLoader(train_set, BATCH_SIZE, True,
    ↳collate_fn=review_collate_fn)
val_loader = DataLoader(val_set, BATCH_SIZE, False,
    ↳collate_fn=review_collate_fn)
test_loader = DataLoader(test_set, BATCH_SIZE, False,
    ↳collate_fn=review_collate_fn)

```

```

[ ]: # Let's print the first batch
batch = next(iter(val_loader))
print(batch)

# (tensor([[ 11,    6,    1, ...,  0,    0,    0],
#          [ 11,  170, 2220, ...,  0,    0,    0],
#          [ 48,    3, 30376, ...,  0,    0,    0],
#          ...,
#          [ 176,   56,   10, ...,  0,    0,    0],
#          [ 239,  534, 1404, ...,  44,  120,    1],

```

```
# [ 2954, 15576, 6, ..., 2678, 65, 1]],
# tensor([0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0,
↪1,
# 0, 0, 0, 1, 1, 1, 0, 0]))

(tensor([[ 11, 6, 1, ..., 0, 0, 0],
[ 11, 170, 2220, ..., 0, 0, 0],
[ 48, 3, 30376, ..., 0, 0, 0],
...,
[ 176, 56, 10, ..., 0, 0, 0],
[ 239, 534, 1404, ..., 44, 120, 1],
[ 2954, 15576, 6, ..., 2678, 65, 1]]), tensor([0, 1, 0, 1,
0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1,
0, 0, 0, 1, 1, 1, 0, 0]))
```

2.3 Task 3: Recurrent Neural Networks

```
[ ]: class SimpleLSTM(nn.Module):
    def __init__(
        self, vocab_size, embedding_dim, hidden_dim, num_layers=1,
↪cell_dropout=0.0
    ):
        """
        Initializes the model by setting up the layers.

        Args:
            vocab_size: number of unique words in the reviews
            embedding_dim: size of the embeddings
            hidden_dim: dimension of the LSTM output
            num_layers: number of LSTM layers
            cell_dropout: dropout applied between the LSTM layers
                        (provide to LSTM constructor as dropout argument)
        """
        super().__init__()

        self.num_layers = num_layers
        self.hidden_dim = hidden_dim

        # YOUR CODE HERE
        self.embedding = nn.Embedding(vocab_size, embedding_dim)

        self.lstm = nn.LSTM(
            embedding_dim, hidden_dim, num_layers=num_layers,
↪dropout=cell_dropout, batch_first=True
        )

        self.fc = nn.Linear(hidden_dim, 1)
```



```

self.sigmoid = nn.Sigmoid()

def forward(self, x):
    """
    Performs a forward pass of the model on some input and hidden state.

    Parameters
    -----
    x: batch as a (batch_size, sequence_length) tensor

    Returns
    -----
    Probability of positive class and the last output of the LSTM.
    """
    # init hidden layer, which is needed for the LSTM
    hidden = self.init_hidden(len(x))

    # YOUR CODE HERE

    # Embedding layer
    embedded = self.embedding(x)

    # LSTM layer
    lstm_out, hidden = self.lstm(embedded, hidden)

    # Take the last output of the LSTM
    last_output = lstm_out[:, -1, :]

    # Linear layer
    logits = self.fc(last_output)

    # Sigmoid activation
    prob = self.sigmoid(logits)

    return prob, last_output

def init_hidden(self, batch_size):
    """Initialize hidden states.

    Returns a tuple of two num_layers x batch_size x hidden_dim tensors
    ↪ (one for
        initial cell states, one for initial hidden states) consisting of all
    ↪ zeros.
    """

```

```

    # YOUR CODE HERE
    # Note: ensure that the returned tensors are located on DEVICE.
    weight = next(self.parameters()).data
    hidden = (
        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_().
        ↪to(DEVICE),
        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_().
        ↪to(DEVICE)
    )
    return hidden

# Test constructor
model = SimpleLSTM(50, 10, 32, 2, 0.1).to(DEVICE)
print(model)

# Should give:
# SimpleLSTM(
#   (embedding): Embedding(50, 10)
#   (lstm): LSTM(10, 32, num_layers=2, batch_first=True, dropout=0.1)
#   (fc): Linear(in_features=32, out_features=1, bias=True)
#   (sigmoid): Sigmoid()
# )

```

```

SimpleLSTM(
  (embedding): Embedding(50, 10)
  (lstm): LSTM(10, 32, num_layers=2, batch_first=True, dropout=0.1)
  (fc): Linear(in_features=32, out_features=1, bias=True)
  (sigmoid): Sigmoid()
)

```

```

[ ]: # Test forward pass
model = SimpleLSTM(50, 10, 32, 2).to(DEVICE)
dummy_data = torch.arange(30, dtype=torch.int, device=DEVICE).reshape(3, 10)

# fix model parameters
for key in model.state_dict():
    model.state_dict()[key][:] = 0.1
probs, states = model(dummy_data)
print(probs)
print(states)

# tensor([[0.9643],
#         [0.9643],
#         [0.9643]], device='cuda:0 or cpu', grad_fn=<SigmoidBackward0>)

```

```

# tensor([[0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
#          [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
#          [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#          0.9985, 0.9985, 0.9985, 0.9985, 0.9985], device='cuda:0 or cpu',
#          grad_fn=<SliceBackward0>))

```

```

tensor([[0.9643],
        [0.9643],
        [0.9643]], grad_fn=<SigmoidBackward0>)
tensor([[0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
        [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
        [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985]], grad_fn=<SliceBackward0>))

```

2.3.1 Task 3d

```

[ ]: @torch.no_grad() # Disables autograd for this function
def reviews_eval(
    model, eval_loader, label="val", loss_fn=torch.nn.functional.
    ↪binary_cross_entropy
):
    model.to(DEVICE)

```

```

model.eval() # sets model to evaluation mode (e.g., relevant for dropout)

total_correct = total_loss = 0
for reviews, labels in eval_loader:
    reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

    # Forward pass: Compute the model's output, reshape it to a vector, and
    # then run the provided loss function.
    # YOUR CODE HERE

    # Receive model output (probabilities)
    output, _ = model(reviews)
    output = output.view(-1)

    # Compute loss
    loss = loss_fn(output, labels.float())

    # Eval stats: Add loss to total_loss and number of correct predictions
    ↪to
    # total_correct.
    # YOUR CODE HERE

    # Add loss to total_loss
    total_loss += loss.item()

    # Add number of correct predictions to total_correct
    total_correct += ((output > 0.5) == labels).sum().item()

print(
    f"
    f"{label} accuracy: {total_correct / len(eval_loader.dataset):.4f}\t"
    f"{label} loss: {total_loss / len(eval_loader):.4f}"
)

```

```

[ ]: # Test your implementation
model = SimpleLSTM(len(dataset.vocab), 10, 10, 1, 0).to(DEVICE)
reviews_eval(model, val_loader)

# Should yield with different but similar numbers:
#

```

```

val accuracy: 0.5225      val loss: 0.6926

```

2.3.2 Task 3e

```
[ ]: def reviews_train(
    model,
    train_loader,
    val_loader,
    lr=0.01,
    epochs=3,
    max_norm=5,
    loss_fn=torch.nn.functional.binary_cross_entropy,
):
    """
    Train a network on the review data

    Args:
        model: Initialized model
        train_loader: Dataloader for the training data
        val_loader: Dataloader for the validation data
        lr: learning rate
        epochs: number of epochs
        max_norm: max norm of gradients for gradient clipping
        loss_fn: Loss function
    """

    # Send the model's parameters to your accelerator (cuda or cpu)
    model = model.to(DEVICE)

    # Define optimizer for the parameters which require gradients (cf. Task 5)
    optimizer = torch.optim.Adam(
        [param for param in model.parameters() if param.requires_grad], lr=lr
    )

    # Let's go
    for epoch in range(epochs):
        total_correct = total_loss = 0
        for reviews, labels in train_loader:
            model.train()

            # Send batch to your accelerator
            reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

            # Forward pass: Compute the model's output, reshape it to a vector,
            ↪ and then
            # run the provided loss function.
            output, _ = model(reviews)
            output = output.view(-1)
            loss = loss_fn(output, labels.float())
```

```

# Backward pass:
# (i) Compute the gradients wrt. the loss
loss.backward()
# (ii) Clip the gradients using
# https://pytorch.org/docs/stable/generated/torch.nn.utils.
↳clip_grad_norm_.html to max_norm
torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm)
# (iii) Run the optimizer
optimizer.step()
# (iv) Clear all accumulated gradients
optimizer.zero_grad()

# Compute epoch statistics:
# (i) Add the loss of this batch to the total_loss
total_loss += loss.item()
# (ii) Add the number of correct predictions (max prob) to
↳total_correct
total_correct += ((output > 0.5) == labels).sum().item()

print(
    f"Epoch {epoch + 1:2}\t"
    f"train accuracy: {total_correct / len(train_loader.dataset):.4f}\t"
    f"train loss: {total_loss / len(train_loader):.4f}"
)

# now validate
reviews_eval(model, val_loader, loss_fn=loss_fn)

```

```

[ ]: # Test your implementation
model = SimpleLSTM(len(dataset.vocab), 10, 10, 1).to(DEVICE)
reviews_train(model, train_loader, val_loader, epochs=5)

# Should yield something like (note: numbers have high variance over runs):
# Epoch 1      train accuracy: 0.4994      train loss: 0.6953
#              val accuracy: 0.4875      val loss: 0.6922
# Epoch 2      train accuracy: 0.5319      train loss: 0.6885
#              val accuracy: 0.5275      val loss: 0.6861
# Epoch 3      train accuracy: 0.6059      train loss: 0.6443
#              val accuracy: 0.5400      val loss: 0.6902
# Epoch 4      train accuracy: 0.6863      train loss: 0.5438
#              val accuracy: 0.5925      val loss: 0.7453
# Epoch 5      train accuracy: 0.8334      train loss: 0.3875
#              val accuracy: 0.7300      val loss: 0.6310

```

```

Epoch 1      train accuracy: 0.5116  train loss: 0.6935
              val accuracy: 0.4850    val loss: 0.6951
Epoch 2      train accuracy: 0.5503  train loss: 0.6814
              val accuracy: 0.5175    val loss: 0.7005

```

Epoch 3	train accuracy: 0.6034	train loss: 0.6346
	val accuracy: 0.5000	val loss: 0.7338
Epoch 4	train accuracy: 0.7141	train loss: 0.5344
	val accuracy: 0.6625	val loss: 0.6996
Epoch 5	train accuracy: 0.8216	train loss: 0.4115
	val accuracy: 0.6650	val loss: 0.7808

2.4 Task 4: Pre-trained Embeddings & Visualization

2.4.1 Task 4b

```
[ ]: # Load Glove embeddings into a plain embedding layer.
vocab = dataset.vocab
glove_embeddings = nn.Embedding(len(vocab), 100, device=DEVICE)
reviews_load_embeddings(glove_embeddings, vocab.get_stoi())
```

Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

```
[ ]: reviews_load_embeddings(glove_embeddings, vocab.get_stoi())
```

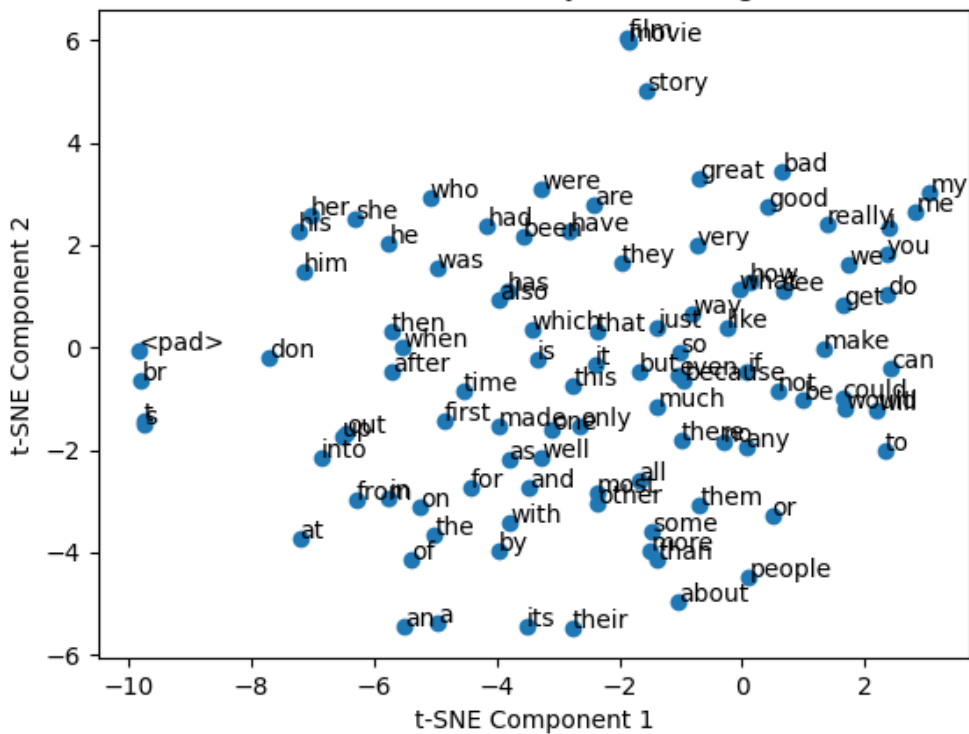
Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

```
[ ]: # Print one embedding
glove_embeddings(torch.tensor(vocab["movie"], device=DEVICE))
```

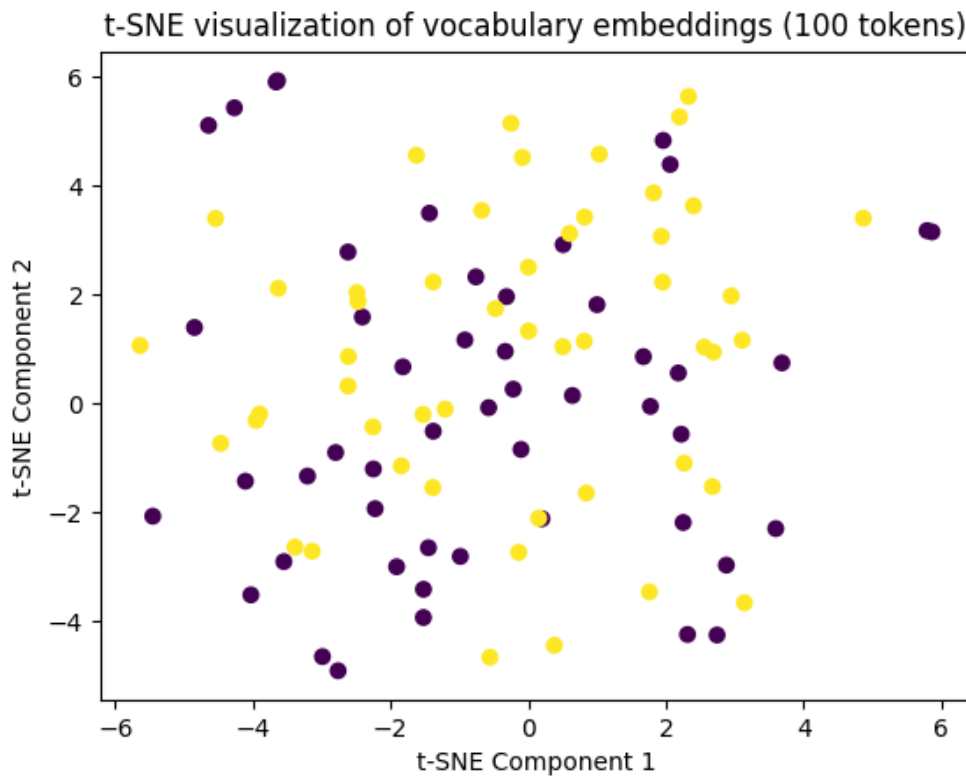
```
[ ]: tensor([ 0.3825,  0.1482,  0.6060, -0.5153,  0.4399,  0.0611, -0.6272, -0.0254,
            0.1643, -0.2210,  0.1442, -0.3721, -0.2168, -0.0890,  0.0979,  0.6561,
            0.6446,  0.4770,  0.8385,  1.6486,  0.8892, -0.1181, -0.0125, -0.5208,
            0.7785,  0.4872, -0.0150, -0.1413, -0.3475, -0.2959,  0.1028,  0.5719,
           -0.0456,  0.0264,  0.5382,  0.3226,  0.4079, -0.0436, -0.1460, -0.4835,
            0.3204,  0.5509, -0.7626,  0.4327,  0.6175, -0.3650, -0.6060, -0.7962,
            0.3929, -0.2367, -0.3472, -0.6120,  0.5475,  0.9481,  0.2094, -2.7771,
           -0.6022,  0.8495,  1.2549,  0.0179, -0.0419,  2.1147, -0.0266, -0.2810,
            0.6812, -0.1417,  0.9925,  0.4988, -0.6754,  0.6417,  0.4230, -0.2791,
            0.0634,  0.6891, -0.3618,  0.0537, -0.1681,  0.1942, -0.4707, -0.1480,
           -0.5899, -0.2797,  0.1679,  0.1057, -1.7601,  0.0088, -0.8333, -0.5836,
           -0.3708, -0.5659,  0.2070,  0.0713,  0.0556, -0.2976, -0.0727, -0.2560,
            0.4269,  0.0589,  0.0911,  0.4728], grad_fn=<EmbeddingBackward0>)
```

```
[ ]: # Plot embeddings of first 100 words using t-SNE
nextplot()
_ = tsne_vocab(glove_embeddings, torch.arange(100), vocab)
```

t-SNE visualization of vocabulary embeddings (100 tokens)



```
[ ]: # You can also specify colors and/or drop the item labels
nextplot()
_ = tsne_vocab(glove_embeddings, torch.arange(100), colors=[0] * 50 + [1] * 50)
```

2.4.2 Adding PCA

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

def PCA_embeddings(embeddings, tokens=None, vocab=None, colors=None,
    ↪ n_components=2):
    """
    Visualize embeddings from an embedding module using PCA.

    Args:
        embeddings: embedding layer of the model
        tokens: limit to the provided embeddings (a tensor of indexes)
        vocab: a vocabulary to label dots in scatter plot with tokens
        colors: colors for each point in the scatter plot
        n_components: number of principal components to use for PCA
    """
    # Convert embeddings to numpy array
    embedding_matrix = embeddings.weight.data.numpy()[ :100]
```

```

if tokens is not None:
    embedding_matrix = embedding_matrix[tokens]
else:
    tokens = torch.arange(len(embedding_matrix))

# Perform PCA
pca = PCA(n_components=n_components)
principal_components = pca.fit_transform(embedding_matrix)


# Plot PCA result
plt.figure(figsize=(8, 6))
plt.scatter(principal_components[:, 0], principal_components[:, 1],
c=colors, alpha=0.5)

# Add token labels to points
if vocab:
    for i in range(len(embedding_matrix)):
        plt.annotate(vocab.get_itos()[tokens[i]], (principal_components[i,
0], principal_components[i, 1]))

plt.title(f'PCA visualization of vocabulary embeddings (100 tokens)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()

PCA_embeddings(glove_embeddings, vocab=vocab, colors="green")

```



2.4.3 Task 4c

```
[ ]: # hyperparameter settings for rest of task 4
vocab_size = len(dataset.vocab)
embedding_dim = 100
hidden_dim = 100
num_layers = 2
n_epochs = 10
cell_dropout = 0.0
```


```
[ ]: # train a plain model
model = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
    ↪ cell_dropout).to(
    DEVICE
)
reviews_train(model, train_loader, val_loader, epochs=n_epochs)

# Should reach a (train) accuracy of >0.9. If not, rerun.
```


```
Epoch 1      train accuracy: 0.4841  train loss: 0.6972
            val accuracy: 0.5000    val loss: 0.6927
```

Epoch 2	train accuracy: 0.5537	train loss: 0.6813
	val accuracy: 0.4650	val loss: 0.7082
Epoch 3	train accuracy: 0.6194	train loss: 0.5905
	val accuracy: 0.5400	val loss: 0.7525
Epoch 4	train accuracy: 0.6759	train loss: 0.4895
	val accuracy: 0.5275	val loss: 0.8901
Epoch 5	train accuracy: 0.7666	train loss: 0.4141
	val accuracy: 0.5675	val loss: 1.0505
Epoch 6	train accuracy: 0.8378	train loss: 0.3466
	val accuracy: 0.5550	val loss: 1.1407
Epoch 7	train accuracy: 0.8819	train loss: 0.2769
	val accuracy: 0.5325	val loss: 1.3425
Epoch 8	train accuracy: 0.9119	train loss: 0.2207
	val accuracy: 0.5550	val loss: 1.3570
Epoch 9	train accuracy: 0.9297	train loss: 0.1852
	val accuracy: 0.5675	val loss: 1.1795
Epoch 10	train accuracy: 0.9475	train loss: 0.1463
	val accuracy: 0.5675	val loss: 1.5681

```
[ ]: # Plot t-SNE embeddings of the thought vectors for training data
      # point color = label
      nextplot()
      _ = tsne_thought(model, train_loader, DEVICE)
```



```
[ ]: # Plot t-SNE embeddings of of the thought vectors for validation data
nextplot()
_ = tsne_thought(model, val_loader, DEVICE)
```



2.4.4 Task 4d

```
[ ]: # Initialize the model with *p*re-trained embeddings with *f*inetuning, then
      ↪ train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
      ↪ cell_dropout)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1 train accuracy: 0.4856 train loss: 0.6947

 val accuracy: 0.4900 val loss: 0.6933

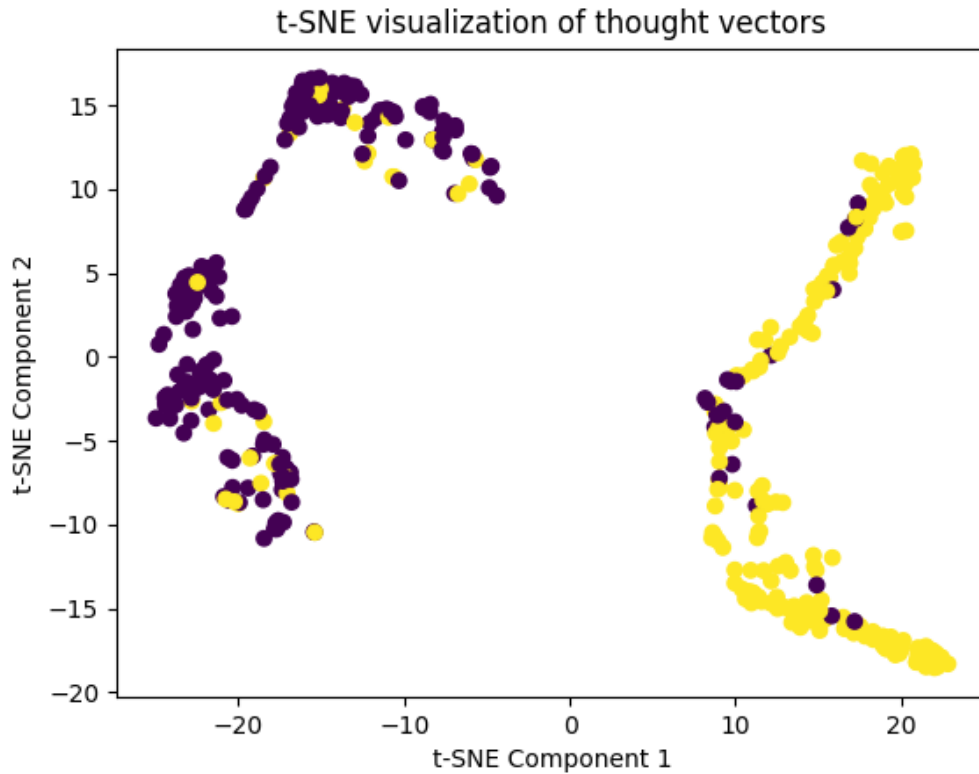
Epoch 2 train accuracy: 0.5416 train loss: 0.6764

	val accuracy: 0.5300	val loss: 0.6979
Epoch 3	train accuracy: 0.7497	train loss: 0.4796
	val accuracy: 0.7050	val loss: 0.5697
Epoch 4	train accuracy: 0.8988	train loss: 0.2700
	val accuracy: 0.8100	val loss: 0.5045
Epoch 5	train accuracy: 0.9594	train loss: 0.1300
	val accuracy: 0.8375	val loss: 0.5846
Epoch 6	train accuracy: 0.9872	train loss: 0.0551
	val accuracy: 0.8075	val loss: 0.6754
Epoch 7	train accuracy: 0.9953	train loss: 0.0214
	val accuracy: 0.8250	val loss: 0.7818
Epoch 8	train accuracy: 0.9978	train loss: 0.0140
	val accuracy: 0.8150	val loss: 0.8342
Epoch 9	train accuracy: 0.9981	train loss: 0.0095
	val accuracy: 0.8225	val loss: 0.8304
Epoch 10	train accuracy: 0.9984	train loss: 0.0053
	val accuracy: 0.8100	val loss: 0.9343

```
[ ]: nextplot()
_ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[ ]: nextplot()
_ = tsne_thought(model_pf, val_loader, DEVICE)
```



2.4.5 Task 4e

```
[ ]: # Initialize the model with *pre-trained embeddings without finetuning, then
      ↪ train
model_p = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
      ↪ cell_dropout)
reviews_load_embeddings(model_p.embedding, vocab.get_stoi())
model_p.embedding.weight.requires_grad = False
reviews_train(model_p, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.4928	train loss: 0.6951
	val accuracy: 0.5125	val loss: 0.6931
Epoch 2	train accuracy: 0.5141	train loss: 0.6929
	val accuracy: 0.5200	val loss: 0.6919
Epoch 3	train accuracy: 0.5178	train loss: 0.6901
	val accuracy: 0.4850	val loss: 0.6932

Epoch 4	train accuracy: 0.5022	train loss: 0.6933
	val accuracy: 0.5175	val loss: 0.6928
Epoch 5	train accuracy: 0.5397	train loss: 0.6808
	val accuracy: 0.5175	val loss: 0.6922
Epoch 6	train accuracy: 0.5709	train loss: 0.6630
	val accuracy: 0.5425	val loss: 0.6798
Epoch 7	train accuracy: 0.5994	train loss: 0.6324
	val accuracy: 0.5225	val loss: 0.7071
Epoch 8	train accuracy: 0.7191	train loss: 0.5594
	val accuracy: 0.6525	val loss: 0.6325
Epoch 9	train accuracy: 0.7428	train loss: 0.5303
	val accuracy: 0.6750	val loss: 0.6083
Epoch 10	train accuracy: 0.7519	train loss: 0.5262
	val accuracy: 0.5825	val loss: 0.6507

```
[ ]: nextplot()
_ = tsne_thought(model_p, train_loader, DEVICE)
```



```
[ ]: nextplot()
_ = tsne_thought(model_p, val_loader, DEVICE)
```

2.5 Task 5: Exploration

2.5.1 Task 5a

Simple RNN class implementation

```
[ ]: # YOUR CODE HERE
class SimpleRNN(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, num_layers=1,
        cell_dropout=0.0, bidirectional=False, cell_type='lstm'):
        super().__init__()

        self.num_layers = num_layers
        self.hidden_dim = hidden_dim
        self.bidirectional = bidirectional
        self.cell_type = cell_type

        self.embedding = nn.Embedding(vocab_size, embedding_dim)

        if self.cell_type == 'lstm':
```

```

        self.rnn = nn.LSTM(embedding_dim, hidden_dim,
↪num_layers=num_layers, dropout=cell_dropout, batch_first=True,
↪bidirectional=bidirectional)
    elif self.cell_type == 'gru':
        self.rnn = nn.GRU(embedding_dim, hidden_dim, num_layers=num_layers,
↪dropout=cell_dropout, batch_first=True, bidirectional=bidirectional)
    else: # Elman
        self.rnn = nn.RNN(embedding_dim, hidden_dim, num_layers=num_layers,
↪dropout=cell_dropout, batch_first=True, bidirectional=bidirectional)

    self.fc = nn.Linear(hidden_dim * 2 if bidirectional else hidden_dim, 1)
    self.sigmoid = nn.Sigmoid()

def forward(self, x):
    hidden = self.init_hidden(len(x))

    embedded = self.embedding(x)
    rnn_out, hidden = self.rnn(embedded, hidden)
    last_output = rnn_out[:, -1, :]

    logits = self.fc(last_output)
    prob = self.sigmoid(logits)

    return prob, last_output

def init_hidden(self, batch_size):
    weight = next(self.parameters()).data
    if self.cell_type == 'lstm':
        hidden = (
            weight.new(self.num_layers * (2 if self.bidirectional else 1),
↪batch_size, self.hidden_dim).zero_().to(DEVICE),
            weight.new(self.num_layers * (2 if self.bidirectional else 1),
↪batch_size, self.hidden_dim).zero_().to(DEVICE)
        )
    else: # Elman RNN and GRU
        hidden = weight.new(self.num_layers * (2 if self.bidirectional else
↪1), batch_size, self.hidden_dim).zero_().to(DEVICE)
    return hidden

```

2.5.2 Train and evaluate the model

Simple RNN is trained and evaluated with the all possible combinations of the following parameters:
- **Dropout:** 0, 0.8 - **RNN cell types:** Elman cell, GRU cell, LSTM cell - **Direction of RNN :** unidirectional, bidirectional.

```

[ ]: # Dropout values
     dropout_values = [0, 0.8]

```

```

# Cell types
cell_types = ['rnn', 'gru', 'lstm']

# Bidirectional values
bidirectional_values = [False, True]

for dropout in dropout_values:
    for cell_type in cell_types:
        for bidirectional in bidirectional_values:
            print(f"Training model with dropout={dropout},  

↪ cell_type={cell_type}, bidirectional={str(bidirectional)}")
            model = SimpleRNN(vocab_size, embedding_dim, hidden_dim,  

↪ num_layers, dropout, bidirectional, cell_type)
            reviews_load_embeddings(model.embedding, vocab.get_stoi())
            reviews_train(model, train_loader, val_loader, epochs=n_epochs)

```

Training model with dropout=0, cell_type=rnn, bidirectional=False
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5009	train loss: 0.7070
	val accuracy: 0.4900	val loss: 0.6977
Epoch 2	train accuracy: 0.5103	train loss: 0.7074
	val accuracy: 0.5225	val loss: 0.7022
Epoch 3	train accuracy: 0.5169	train loss: 0.7050
	val accuracy: 0.5175	val loss: 0.7030
Epoch 4	train accuracy: 0.5487	train loss: 0.6772
	val accuracy: 0.5300	val loss: 0.7158
Epoch 5	train accuracy: 0.6028	train loss: 0.6299
	val accuracy: 0.4775	val loss: 0.7776
Epoch 6	train accuracy: 0.6278	train loss: 0.5906
	val accuracy: 0.5075	val loss: 0.7360
Epoch 7	train accuracy: 0.6494	train loss: 0.5582
	val accuracy: 0.5400	val loss: 0.7344
Epoch 8	train accuracy: 0.6434	train loss: 0.5669
	val accuracy: 0.5425	val loss: 0.8228
Epoch 9	train accuracy: 0.6669	train loss: 0.5287
	val accuracy: 0.5375	val loss: 0.8146
Epoch 10	train accuracy: 0.6725	train loss: 0.5062
	val accuracy: 0.5000	val loss: 0.9907

Training model with dropout=0, cell_type=rnn, bidirectional=True
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4913	train loss: 0.7328
	val accuracy: 0.4900	val loss: 0.7192
Epoch 2	train accuracy: 0.5194	train loss: 0.7181
	val accuracy: 0.5025	val loss: 0.7603

Epoch 3	train accuracy: 0.5553	train loss: 0.6893
	val accuracy: 0.5100	val loss: 0.7046
Epoch 4	train accuracy: 0.5691	train loss: 0.6802
	val accuracy: 0.5100	val loss: 0.7298
Epoch 5	train accuracy: 0.5956	train loss: 0.6530
	val accuracy: 0.5350	val loss: 0.6974
Epoch 6	train accuracy: 0.5863	train loss: 0.6386
	val accuracy: 0.4950	val loss: 0.7366
Epoch 7	train accuracy: 0.6209	train loss: 0.6276
	val accuracy: 0.5275	val loss: 0.7529
Epoch 8	train accuracy: 0.6059	train loss: 0.6213
	val accuracy: 0.5150	val loss: 0.7597
Epoch 9	train accuracy: 0.6047	train loss: 0.6424
	val accuracy: 0.5275	val loss: 0.7577
Epoch 10	train accuracy: 0.6094	train loss: 0.6001
	val accuracy: 0.4775	val loss: 0.7557

Training model with dropout=0, cell_type=gru, bidirectional=False
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5828	train loss: 0.6696
	val accuracy: 0.7775	val loss: 0.4800
Epoch 2	train accuracy: 0.9053	train loss: 0.2511
	val accuracy: 0.8650	val loss: 0.2973
Epoch 3	train accuracy: 0.9944	train loss: 0.0242
	val accuracy: 0.8125	val loss: 0.7040
Epoch 4	train accuracy: 0.9984	train loss: 0.0058
	val accuracy: 0.8575	val loss: 0.6349
Epoch 5	train accuracy: 1.0000	train loss: 0.0003
	val accuracy: 0.8600	val loss: 0.7511
Epoch 6	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8600	val loss: 0.7970
Epoch 7	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8575	val loss: 0.8287
Epoch 8	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8575	val loss: 0.8556
Epoch 9	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8575	val loss: 0.8788
Epoch 10	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8575	val loss: 0.8993

Training model with dropout=0, cell_type=gru, bidirectional=True
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4988	train loss: 0.7147
	val accuracy: 0.4800	val loss: 0.7104
Epoch 2	train accuracy: 0.7197	train loss: 0.5031
	val accuracy: 0.8350	val loss: 0.4141
Epoch 3	train accuracy: 0.9681	train loss: 0.0965
	val accuracy: 0.8575	val loss: 0.4100

Epoch 4	train accuracy: 0.9969	train loss: 0.0125
	val accuracy: 0.8050	val loss: 0.6617
Epoch 5	train accuracy: 0.9991	train loss: 0.0019
	val accuracy: 0.8350	val loss: 0.9350
Epoch 6	train accuracy: 0.9997	train loss: 0.0005
	val accuracy: 0.8475	val loss: 0.8510
Epoch 7	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8550	val loss: 0.9483
Epoch 8	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8525	val loss: 0.9755
Epoch 9	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8525	val loss: 0.9997
Epoch 10	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8525	val loss: 1.0208

Training model with dropout=0, cell_type=lstm, bidirectional=False
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5012	train loss: 0.6943
	val accuracy: 0.5050	val loss: 0.6938
Epoch 2	train accuracy: 0.5337	train loss: 0.6898
	val accuracy: 0.4975	val loss: 0.6982
Epoch 3	train accuracy: 0.5934	train loss: 0.6467
	val accuracy: 0.5000	val loss: 0.7164
Epoch 4	train accuracy: 0.6931	train loss: 0.5581
	val accuracy: 0.7450	val loss: 0.5832
Epoch 5	train accuracy: 0.8628	train loss: 0.3542
	val accuracy: 0.7725	val loss: 0.5599
Epoch 6	train accuracy: 0.9534	train loss: 0.1525
	val accuracy: 0.7350	val loss: 0.7497
Epoch 7	train accuracy: 0.9791	train loss: 0.0654
	val accuracy: 0.7950	val loss: 0.6579
Epoch 8	train accuracy: 0.9906	train loss: 0.0383
	val accuracy: 0.7750	val loss: 0.7919
Epoch 9	train accuracy: 0.9950	train loss: 0.0209
	val accuracy: 0.7475	val loss: 0.8680
Epoch 10	train accuracy: 0.9947	train loss: 0.0131
	val accuracy: 0.7725	val loss: 0.9023

Training model with dropout=0, cell_type=lstm, bidirectional=True
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5006	train loss: 0.7016
	val accuracy: 0.4900	val loss: 0.6972
Epoch 2	train accuracy: 0.5181	train loss: 0.6920
	val accuracy: 0.4675	val loss: 0.7184
Epoch 3	train accuracy: 0.5559	train loss: 0.6782
	val accuracy: 0.5075	val loss: 0.7111
Epoch 4	train accuracy: 0.5750	train loss: 0.6254
	val accuracy: 0.5050	val loss: 0.7767

Epoch 5	train accuracy: 0.6016	train loss: 0.5927
	val accuracy: 0.4850	val loss: 0.8704
Epoch 6	train accuracy: 0.7056	train loss: 0.5238
	val accuracy: 0.5900	val loss: 0.9097
Epoch 7	train accuracy: 0.8125	train loss: 0.4097
	val accuracy: 0.6375	val loss: 0.7661
Epoch 8	train accuracy: 0.8878	train loss: 0.2655
	val accuracy: 0.7075	val loss: 0.7162
Epoch 9	train accuracy: 0.9494	train loss: 0.1424
	val accuracy: 0.7900	val loss: 0.6476
Epoch 10	train accuracy: 0.9812	train loss: 0.0601
	val accuracy: 0.7750	val loss: 0.8637

Training model with dropout=0.8, cell_type=rnn, bidirectional=False
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5141	train loss: 0.7134
	val accuracy: 0.4875	val loss: 0.7439
Epoch 2	train accuracy: 0.4791	train loss: 0.7135
	val accuracy: 0.5125	val loss: 0.7524
Epoch 3	train accuracy: 0.4891	train loss: 0.7167
	val accuracy: 0.5275	val loss: 0.6945
Epoch 4	train accuracy: 0.4947	train loss: 0.7103
	val accuracy: 0.5225	val loss: 0.6971
Epoch 5	train accuracy: 0.4844	train loss: 0.7118
	val accuracy: 0.5200	val loss: 0.6984
Epoch 6	train accuracy: 0.5081	train loss: 0.7048
	val accuracy: 0.5125	val loss: 0.7070
Epoch 7	train accuracy: 0.4859	train loss: 0.7154
	val accuracy: 0.5200	val loss: 0.7059
Epoch 8	train accuracy: 0.5234	train loss: 0.7087
	val accuracy: 0.4875	val loss: 0.7109
Epoch 9	train accuracy: 0.4856	train loss: 0.7153
	val accuracy: 0.5225	val loss: 0.6923
Epoch 10	train accuracy: 0.5022	train loss: 0.7036
	val accuracy: 0.5225	val loss: 0.6941

Training model with dropout=0.8, cell_type=rnn, bidirectional=True
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5050	train loss: 0.7461
	val accuracy: 0.4925	val loss: 0.7155
Epoch 2	train accuracy: 0.5081	train loss: 0.7196
	val accuracy: 0.4925	val loss: 0.7532
Epoch 3	train accuracy: 0.4941	train loss: 0.7296
	val accuracy: 0.5200	val loss: 0.8333
Epoch 4	train accuracy: 0.5288	train loss: 0.7074
	val accuracy: 0.4950	val loss: 0.7150
Epoch 5	train accuracy: 0.5553	train loss: 0.6820
	val accuracy: 0.5325	val loss: 0.7119

Epoch 6	train accuracy: 0.5925	train loss: 0.6562
	val accuracy: 0.5000	val loss: 0.7358
Epoch 7	train accuracy: 0.6062	train loss: 0.6292
	val accuracy: 0.5100	val loss: 0.8313
Epoch 8	train accuracy: 0.6034	train loss: 0.6463
	val accuracy: 0.5275	val loss: 0.8154
Epoch 9	train accuracy: 0.5984	train loss: 0.6204
	val accuracy: 0.5250	val loss: 0.7760
Epoch 10	train accuracy: 0.6069	train loss: 0.6278
	val accuracy: 0.4925	val loss: 0.7899

Training model with dropout=0.8, cell_type=gru, bidirectional=False

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5031	train loss: 0.7122
	val accuracy: 0.4825	val loss: 0.7708
Epoch 2	train accuracy: 0.6175	train loss: 0.6262
	val accuracy: 0.8000	val loss: 0.4122
Epoch 3	train accuracy: 0.9303	train loss: 0.1884
	val accuracy: 0.8600	val loss: 0.3518
Epoch 4	train accuracy: 0.9847	train loss: 0.0484
	val accuracy: 0.8300	val loss: 0.4822
Epoch 5	train accuracy: 0.9947	train loss: 0.0174
	val accuracy: 0.8450	val loss: 0.5877
Epoch 6	train accuracy: 0.9972	train loss: 0.0099
	val accuracy: 0.8575	val loss: 0.6636
Epoch 7	train accuracy: 0.9984	train loss: 0.0029
	val accuracy: 0.8325	val loss: 1.0184
Epoch 8	train accuracy: 0.9978	train loss: 0.0054
	val accuracy: 0.8550	val loss: 0.8904
Epoch 9	train accuracy: 0.9931	train loss: 0.0206
	val accuracy: 0.8450	val loss: 0.8110
Epoch 10	train accuracy: 0.9956	train loss: 0.0136
	val accuracy: 0.8475	val loss: 0.7808

Training model with dropout=0.8, cell_type=gru, bidirectional=True

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5056	train loss: 0.7162
	val accuracy: 0.4900	val loss: 0.7042
Epoch 2	train accuracy: 0.5309	train loss: 0.7044
	val accuracy: 0.5800	val loss: 0.6540
Epoch 3	train accuracy: 0.8516	train loss: 0.3546
	val accuracy: 0.8075	val loss: 0.4432
Epoch 4	train accuracy: 0.9781	train loss: 0.0676
	val accuracy: 0.8550	val loss: 0.3666
Epoch 5	train accuracy: 0.9925	train loss: 0.0308
	val accuracy: 0.8725	val loss: 0.4024
Epoch 6	train accuracy: 0.9950	train loss: 0.0179
	val accuracy: 0.8650	val loss: 0.4109

```

Epoch 7      train accuracy: 0.9953  train loss: 0.0220
              val accuracy: 0.8875    val loss: 0.3903
Epoch 8      train accuracy: 0.9962  train loss: 0.0132
              val accuracy: 0.8825    val loss: 0.3867
Epoch 9      train accuracy: 0.9969  train loss: 0.0115
              val accuracy: 0.8825    val loss: 0.4477
Epoch 10     train accuracy: 0.9969  train loss: 0.0098
              val accuracy: 0.8875    val loss: 0.5010
Training model with dropout=0.8, cell_type=lstm, bidirectional=False
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
Epoch 1      train accuracy: 0.5000  train loss: 0.6980
              val accuracy: 0.4950    val loss: 0.6956
Epoch 2      train accuracy: 0.4894  train loss: 0.6950
              val accuracy: 0.5200    val loss: 0.6935
Epoch 3      train accuracy: 0.5047  train loss: 0.6944
              val accuracy: 0.4900    val loss: 0.6941
Epoch 4      train accuracy: 0.5138  train loss: 0.6939
              val accuracy: 0.4900    val loss: 0.6954
Epoch 5      train accuracy: 0.4959  train loss: 0.6949
              val accuracy: 0.4800    val loss: 0.6950
Epoch 6      train accuracy: 0.5084  train loss: 0.6953
              val accuracy: 0.5000    val loss: 0.6964
Epoch 7      train accuracy: 0.5503  train loss: 0.6843
              val accuracy: 0.5075    val loss: 0.7025
Epoch 8      train accuracy: 0.5906  train loss: 0.6552
              val accuracy: 0.5650    val loss: 0.6960
Epoch 9      train accuracy: 0.6228  train loss: 0.6114
              val accuracy: 0.5850    val loss: 0.7594
Epoch 10     train accuracy: 0.6416  train loss: 0.5660
              val accuracy: 0.5275    val loss: 0.8274
Training model with dropout=0.8, cell_type=lstm, bidirectional=True
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
Epoch 1      train accuracy: 0.5066  train loss: 0.6989
              val accuracy: 0.4950    val loss: 0.6955
Epoch 2      train accuracy: 0.5050  train loss: 0.6947
              val accuracy: 0.5025    val loss: 0.6965
Epoch 3      train accuracy: 0.5309  train loss: 0.6899
              val accuracy: 0.5600    val loss: 0.6840
Epoch 4      train accuracy: 0.5856  train loss: 0.6422
              val accuracy: 0.4875    val loss: 0.7429
Epoch 5      train accuracy: 0.6234  train loss: 0.6057
              val accuracy: 0.5400    val loss: 0.7843
Epoch 6      train accuracy: 0.6575  train loss: 0.5550
              val accuracy: 0.5625    val loss: 0.7665
Epoch 7      train accuracy: 0.7472  train loss: 0.4787
              val accuracy: 0.5800    val loss: 0.7580

```


Epoch 8	train accuracy: 0.7844	train loss: 0.4494
	val accuracy: 0.6875	val loss: 0.8497
Epoch 9	train accuracy: 0.8550	train loss: 0.3500
	val accuracy: 0.7375	val loss: 0.6093
Epoch 10	train accuracy: 0.8741	train loss: 0.3255
	val accuracy: 0.6350	val loss: 0.8643

Simple RNN is trained and evaluated with the all possible combinations of the following parameters:
- **Dropout:** 0.4 - **RNN cell types:** Elman cell, GRU cell, LSTM cell - **Direction of RNN :** unidirectional, bidirectional.

```
[ ]: # Dropout values
dropout_values = [0.4]

# Cell types
cell_types = ['rnn', 'gru', 'lstm']

# Bidirectional values
bidirectional_values = [False, True]

for dropout in dropout_values:
    for cell_type in cell_types:
        for bidirectional in bidirectional_values:
            print(f"Training model with dropout={dropout},
↪cell_type={cell_type}, bidirectional={str(bidirectional)}")
            model = SimpleRNN(vocab_size, embedding_dim, hidden_dim,
↪num_layers, dropout, bidirectional, cell_type)
            reviews_load_embeddings(model.embedding, vocab.get_stoi())
            reviews_train(model, train_loader, val_loader, epochs=n_epochs)
```

Training model with dropout=0.4, cell_type=rnn, bidirectional=False
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4884	train loss: 0.7157
	val accuracy: 0.5175	val loss: 0.6949
Epoch 2	train accuracy: 0.5044	train loss: 0.7055
	val accuracy: 0.4850	val loss: 0.7427
Epoch 3	train accuracy: 0.5044	train loss: 0.7066
	val accuracy: 0.5200	val loss: 0.6956
Epoch 4	train accuracy: 0.5012	train loss: 0.7095
	val accuracy: 0.4900	val loss: 0.6954
Epoch 5	train accuracy: 0.5269	train loss: 0.6979
	val accuracy: 0.5200	val loss: 0.7268
Epoch 6	train accuracy: 0.5875	train loss: 0.6570
	val accuracy: 0.5000	val loss: 0.7310
Epoch 7	train accuracy: 0.6378	train loss: 0.5750
	val accuracy: 0.5475	val loss: 0.8477
Epoch 8	train accuracy: 0.6663	train loss: 0.5276

	val accuracy: 0.5450	val loss: 0.7639
Epoch 9	train accuracy: 0.6847	train loss: 0.4862
	val accuracy: 0.5275	val loss: 0.9407
Epoch 10	train accuracy: 0.6891	train loss: 0.4663
	val accuracy: 0.5025	val loss: 0.9573

Training model with dropout=0.4, cell_type=rnn, bidirectional=True
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5009	train loss: 0.7402
	val accuracy: 0.5175	val loss: 0.6968
Epoch 2	train accuracy: 0.5319	train loss: 0.7139
	val accuracy: 0.5200	val loss: 0.7179
Epoch 3	train accuracy: 0.5500	train loss: 0.6861
	val accuracy: 0.4850	val loss: 0.7048
Epoch 4	train accuracy: 0.5887	train loss: 0.6619
	val accuracy: 0.5300	val loss: 0.7248
Epoch 5	train accuracy: 0.6181	train loss: 0.6547
	val accuracy: 0.4975	val loss: 0.7630
Epoch 6	train accuracy: 0.6275	train loss: 0.5774
	val accuracy: 0.4925	val loss: 0.8133
Epoch 7	train accuracy: 0.6516	train loss: 0.5712
	val accuracy: 0.4875	val loss: 0.8403
Epoch 8	train accuracy: 0.6647	train loss: 0.5330
	val accuracy: 0.5050	val loss: 0.8799
Epoch 9	train accuracy: 0.6597	train loss: 0.5203
	val accuracy: 0.4975	val loss: 0.9422
Epoch 10	train accuracy: 0.6809	train loss: 0.5035
	val accuracy: 0.4975	val loss: 0.9562

Training model with dropout=0.4, cell_type=gru, bidirectional=False
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5384	train loss: 0.6938
	val accuracy: 0.6375	val loss: 0.6885
Epoch 2	train accuracy: 0.8878	train loss: 0.2898
	val accuracy: 0.9050	val loss: 0.2916
Epoch 3	train accuracy: 0.9941	train loss: 0.0238
	val accuracy: 0.8825	val loss: 0.4388
Epoch 4	train accuracy: 0.9991	train loss: 0.0049
	val accuracy: 0.8800	val loss: 0.5274
Epoch 5	train accuracy: 1.0000	train loss: 0.0004
	val accuracy: 0.8575	val loss: 0.7591
Epoch 6	train accuracy: 0.9988	train loss: 0.0038
	val accuracy: 0.8525	val loss: 0.6040
Epoch 7	train accuracy: 0.9994	train loss: 0.0026
	val accuracy: 0.8650	val loss: 0.8175
Epoch 8	train accuracy: 0.9931	train loss: 0.0201
	val accuracy: 0.8575	val loss: 0.7603
Epoch 9	train accuracy: 0.9975	train loss: 0.0067

	val accuracy: 0.8475	val loss: 0.7740
Epoch 10	train accuracy: 0.9981	train loss: 0.0065
	val accuracy: 0.7950	val loss: 1.1038

Training model with dropout=0.4, cell_type=gru, bidirectional=True
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5784	train loss: 0.6623
	val accuracy: 0.7625	val loss: 0.4825
Epoch 2	train accuracy: 0.8956	train loss: 0.2631
	val accuracy: 0.8300	val loss: 0.3866
Epoch 3	train accuracy: 0.9891	train loss: 0.0357
	val accuracy: 0.8450	val loss: 0.6082
Epoch 4	train accuracy: 0.9981	train loss: 0.0051
	val accuracy: 0.8325	val loss: 0.8242
Epoch 5	train accuracy: 1.0000	train loss: 0.0004
	val accuracy: 0.8050	val loss: 1.1303
Epoch 6	train accuracy: 1.0000	train loss: 0.0002
	val accuracy: 0.8250	val loss: 1.0451
Epoch 7	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8250	val loss: 1.0813
Epoch 8	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8250	val loss: 1.1199
Epoch 9	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8275	val loss: 1.1513
Epoch 10	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8275	val loss: 1.1752

Training model with dropout=0.4, cell_type=lstm, bidirectional=False
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4844	train loss: 0.6967
	val accuracy: 0.4900	val loss: 0.6932
Epoch 2	train accuracy: 0.5062	train loss: 0.6936
	val accuracy: 0.4825	val loss: 0.6946
Epoch 3	train accuracy: 0.5522	train loss: 0.6791
	val accuracy: 0.4975	val loss: 0.7008
Epoch 4	train accuracy: 0.6166	train loss: 0.6082
	val accuracy: 0.5200	val loss: 0.7357
Epoch 5	train accuracy: 0.6803	train loss: 0.5060
	val accuracy: 0.7425	val loss: 0.7125
Epoch 6	train accuracy: 0.8625	train loss: 0.3316
	val accuracy: 0.7950	val loss: 0.6042
Epoch 7	train accuracy: 0.9394	train loss: 0.1897
	val accuracy: 0.8200	val loss: 0.6369
Epoch 8	train accuracy: 0.9709	train loss: 0.1035
	val accuracy: 0.8500	val loss: 0.6389
Epoch 9	train accuracy: 0.9825	train loss: 0.0706
	val accuracy: 0.8550	val loss: 0.7070
Epoch 10	train accuracy: 0.9897	train loss: 0.0469

```

                val accuracy: 0.8575          val loss: 0.6696
Training model with dropout=0.4, cell_type=lstm, bidirectional=True
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
Epoch 1      train accuracy: 0.4991  train loss: 0.7016
              val accuracy: 0.4875    val loss: 0.6948
Epoch 2      train accuracy: 0.5409  train loss: 0.6872
              val accuracy: 0.4775    val loss: 0.6972
Epoch 3      train accuracy: 0.5922  train loss: 0.6517
              val accuracy: 0.4825    val loss: 0.7022
Epoch 4      train accuracy: 0.6069  train loss: 0.6123
              val accuracy: 0.5375    val loss: 0.8028
Epoch 5      train accuracy: 0.6706  train loss: 0.5238
              val accuracy: 0.7400    val loss: 0.6853
Epoch 6      train accuracy: 0.8803  train loss: 0.3100
              val accuracy: 0.8025    val loss: 0.6851
Epoch 7      train accuracy: 0.9375  train loss: 0.1797
              val accuracy: 0.7975    val loss: 0.5517
Epoch 8      train accuracy: 0.9728  train loss: 0.0963
              val accuracy: 0.7300    val loss: 0.9135
Epoch 9      train accuracy: 0.9803  train loss: 0.0744
              val accuracy: 0.8350    val loss: 0.6508
Epoch 10     train accuracy: 0.9912  train loss: 0.0353
              val accuracy: 0.8125    val loss: 0.6955

```

2.5.3 Task 5b

According to our research, the model with the highest performance (val accuracy: 0.9050, val loss: 0.2916) has the following parameters:

- **Dropout:** 0.4
- **RNN cell type:** GRU cell
- **Direction of RNN :** unidirectional

P.s. the highest performance has been reached on the 2nd Epoch of training process.

2.5.4 Optimization of the best model

Here we fix aforementioned parameters of the model and try to discover the best set of hyper parameters for it

```

[ ]: # hyperparameter settings
hidden_dims = [50, 100]
num_layers_values = [1, 2]
n_epochs_values = [5]

for hidden_dim in hidden_dims:
    for num_layers in num_layers_values:
        for n_epochs in n_epochs_values:

```

```

print(f"Training model with hidden_dim={hidden_dim},
num_layers={num_layers}, n_epochs={n_epochs}")
model = SimpleRNN(vocab_size, embedding_dim, hidden_dim,
num_layers, 0.4, False, 'gru')
reviews_load_embeddings(model.embedding, vocab.get_stoi())
reviews_train(model, train_loader, val_loader, epochs=n_epochs)

```

Training model with hidden_dim=50, num_layers=1, n_epochs=5

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.6428	train loss: 0.6044
	val accuracy: 0.8550	val loss: 0.3222
Epoch 2	train accuracy: 0.9519	train loss: 0.1377
	val accuracy: 0.8775	val loss: 0.3597
Epoch 3	train accuracy: 0.9981	train loss: 0.0090
	val accuracy: 0.8700	val loss: 0.5382
Epoch 4	train accuracy: 1.0000	train loss: 0.0016
	val accuracy: 0.8775	val loss: 0.5922
Epoch 5	train accuracy: 1.0000	train loss: 0.0008
	val accuracy: 0.8800	val loss: 0.6587

Training model with hidden_dim=50, num_layers=2, n_epochs=5

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5837	train loss: 0.6547
	val accuracy: 0.8300	val loss: 0.4110
Epoch 2	train accuracy: 0.9203	train loss: 0.2124
	val accuracy: 0.8425	val loss: 0.3869
Epoch 3	train accuracy: 0.9922	train loss: 0.0223
	val accuracy: 0.8525	val loss: 0.5517
Epoch 4	train accuracy: 0.9984	train loss: 0.0053
	val accuracy: 0.8700	val loss: 0.6525
Epoch 5	train accuracy: 1.0000	train loss: 0.0003
	val accuracy: 0.8750	val loss: 0.7324

Training model with hidden_dim=100, num_layers=1, n_epochs=5

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5600	train loss: 0.6861
	val accuracy: 0.7900	val loss: 0.5036
Epoch 2	train accuracy: 0.9103	train loss: 0.2391
	val accuracy: 0.8775	val loss: 0.3243
Epoch 3	train accuracy: 0.9888	train loss: 0.0382
	val accuracy: 0.8425	val loss: 0.5253
Epoch 4	train accuracy: 0.9984	train loss: 0.0102
	val accuracy: 0.8775	val loss: 0.4711
Epoch 5	train accuracy: 0.9997	train loss: 0.0030
	val accuracy: 0.8625	val loss: 0.5473

Training model with hidden_dim=100, num_layers=2, n_epochs=5

Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5112	train loss: 0.7109
	val accuracy: 0.4925	val loss: 0.6999
Epoch 2	train accuracy: 0.7419	train loss: 0.4654
	val accuracy: 0.8775	val loss: 0.3355
Epoch 3	train accuracy: 0.9794	train loss: 0.0725
	val accuracy: 0.8850	val loss: 0.3818
Epoch 4	train accuracy: 0.9978	train loss: 0.0088
	val accuracy: 0.8525	val loss: 0.5200
Epoch 5	train accuracy: 1.0000	train loss: 0.0013
	val accuracy: 0.8900	val loss: 0.5719

Visualization of Model's Validation Accuracy and Loss Based on the Epoch Hyperparameter These visualizations reveal that the model achieves its peak validation accuracy and minimum validation loss during the second Epoch.

```
[ ]: import matplotlib.pyplot as plt

# lists are returned from training function
# For simplicity, let's only consider validation accuracy here
val_accuracies_50_1_5 = [0.8550, 0.8775, 0.8700, 0.8775, 0.8800]
val_accuracies_50_2_5 = [0.8300, 0.8425, 0.8525, 0.8700, 0.8750]
val_accuracies_100_1_5 = [0.7900, 0.8775, 0.8425, 0.8775, 0.8625]
val_accuracies_100_2_5 = [0.4925, 0.8775, 0.8850, 0.8525, 0.8900]
epochs = range(1, len(val_accuracies_50_1_5) + 1)

# Plot validation accuracy
plt.figure(figsize=(12, 6))
plt.plot(epochs, val_accuracies_50_1_5, 'b', label='hidden_dim=50, num_layers=1, n_epochs=5')
plt.plot(epochs, val_accuracies_50_2_5, 'r', label='hidden_dim=50, num_layers=2, n_epochs=5')
plt.plot(epochs, val_accuracies_100_1_5, 'g', label='hidden_dim=100, num_layers=1, n_epochs=5')
plt.plot(epochs, val_accuracies_100_2_5, 'y', label='hidden_dim=100, num_layers=2, n_epochs=5')
plt.title('Validation accuracy over epochs for different hyperparameters')
plt.xlabel('Epochs')
plt.ylabel('Validation accuracy')
plt.legend()
plt.show()


# lists are returned from training function
val_losses_50_1_5 = [0.3222, 0.3597, 0.5382, 0.5922, 0.6587]
val_losses_50_2_5 = [0.4110, 0.3869, 0.5517, 0.6525, 0.7324]
val_losses_100_1_5 = [0.5036, 0.3243, 0.5253, 0.4711, 0.5473]
```

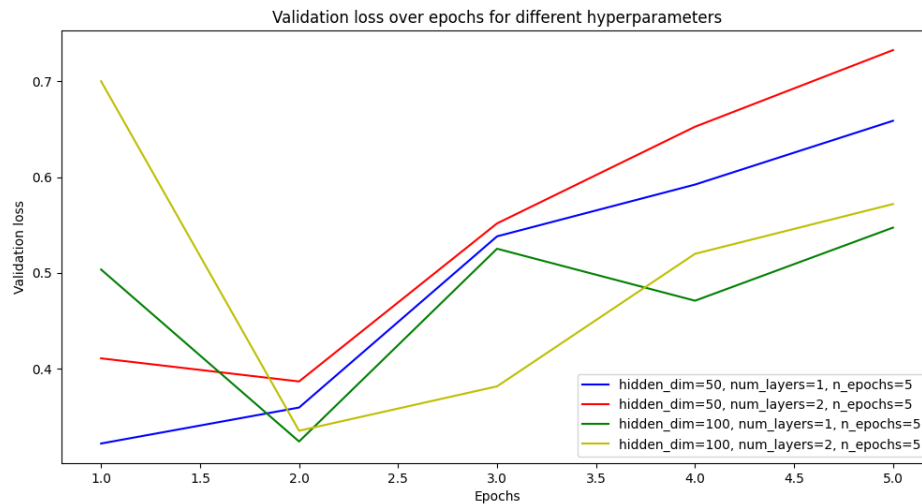
```

val_losses_100_2_5 = [0.6999, 0.3355, 0.3818, 0.5200, 0.5719]

# Plot validation loss
plt.figure(figsize=(12, 6))
plt.plot(epochs, val_losses_50_1_5, 'b', label='hidden_dim=50, num_layers=1, n_epochs=5')
plt.plot(epochs, val_losses_50_2_5, 'r', label='hidden_dim=50, num_layers=2, n_epochs=5')
plt.plot(epochs, val_losses_100_1_5, 'g', label='hidden_dim=100, num_layers=1, n_epochs=5')
plt.plot(epochs, val_losses_100_2_5, 'y', label='hidden_dim=100, num_layers=2, n_epochs=5')
plt.title('Validation loss over epochs for different hyperparameters')
plt.xlabel('Epochs')
plt.ylabel('Validation loss')
plt.legend()
plt.show()

```





Visualization of Model's Validation Accuracy and Loss from the Second Epoch (Peak Performance) The plot reveals two optimal sets of hyperparameters:: 1. hidden_dim=100, num_layers=1: accuracy = 0.8775, loss = 0.3243 2. hidden_dim=100, num_layers=2: accuracy = 0.8775, loss = 0.3355

Given that the accuracy for both sets is identical, and the loss values are only marginally different, we need to consider other factors. We have a moderate n_epoch=2 value and a relatively high dropout=0.4 parameter, which introduces a higher level of regularization. Taking these into account, **we opt for the second set of hyperparameters (hidden_dim=100, num_layers=2)**. This choice enables the second layer of GRUs to detect more complex data patterns, counterbalancing the level of regularization introduced by the n_epoch and dropout hyperparameters.

```
[ ]: import numpy as np

# Validation accuracy and loss at the second epoch for different hyperparameters
val_accuracies = [0.8775, 0.8425, 0.8775, 0.8775]
val_losses = [0.3597, 0.3869, 0.3243, 0.3355]
labels = ['(50, 1)', '(50, 2)', '(100, 1)', '(100, 2)']

x = np.arange(len(labels)) # the label locations
width = 0.35 # the width of the bars

fig, ax1 = plt.subplots()

# Plot validation accuracy
rects1 = ax1.bar(x - width/2, val_accuracies, width, label='Validation_
↳Accuracy')
```



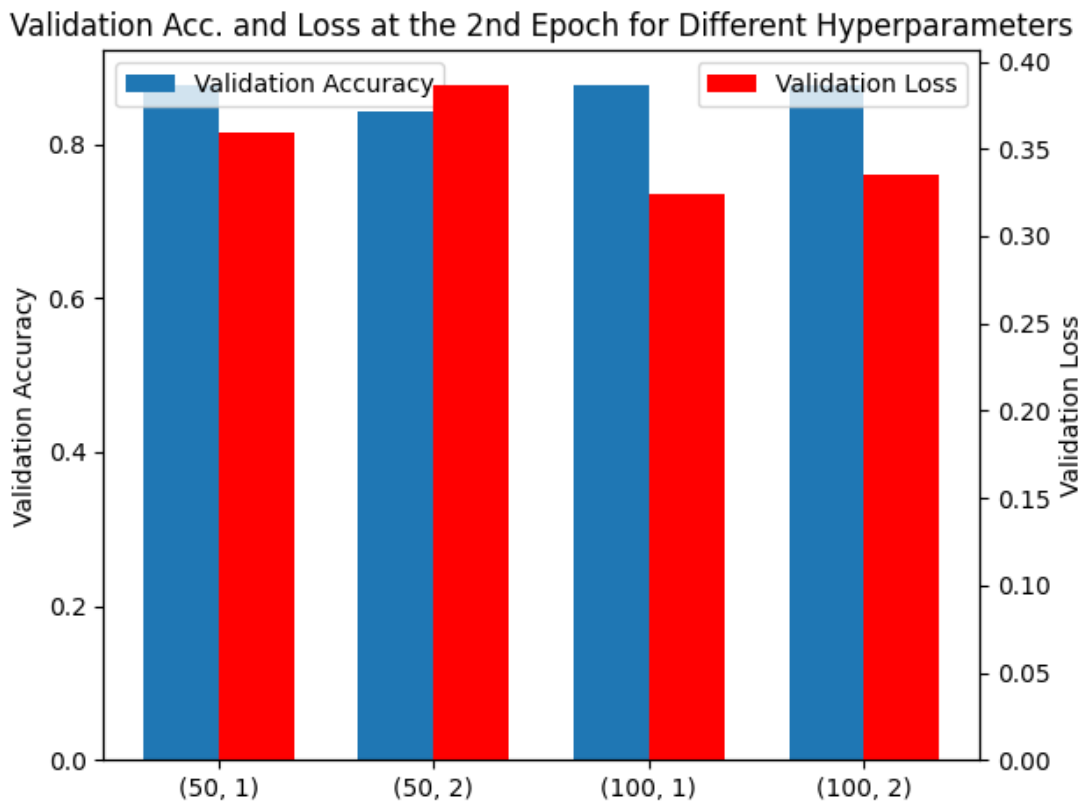
```

ax1.set_ylabel('Validation Accuracy')
ax1.set_title('Validation Acc. and Loss at the 2nd Epoch for Different_
↳Hyperparameters')
ax1.set_xticks(x)
ax1.set_xticklabels(labels)
ax1.legend(loc='upper left')

ax2 = ax1.twinx()
# Plot validation loss
rects2 = ax2.bar(x + width/2, val_losses, width, label='Validation Loss',
↳color='r')
ax2.set_ylabel('Validation Loss')
ax2.legend(loc='upper right')

fig.tight_layout()
plt.show()

```



2.5.5 Evaluation of the optimized model

```
[ ]: # Best hyperparameters
hidden_dim = 100
num_layers = 2
n_epochs = 2
dropout = 0.4
bidirectional = False
cell_type = 'gru'

print(f"Training model with hidden_dim={hidden_dim}, num_layers={num_layers},
      ↪n_epochs={n_epochs}")
model = SimpleRNN(vocab_size, embedding_dim, hidden_dim, num_layers, dropout,
                  ↪bidirectional, cell_type)
reviews_load_embeddings(model.embedding, vocab.get_stoi())
reviews_train(model, train_loader, val_loader, epochs=n_epochs)

reviews_eval(model, test_loader, label="test")
```

```
Training model with hidden_dim=100, num_layers=2, n_epochs=2
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
Epoch 1      train accuracy: 0.5900  train loss: 0.6332
            val accuracy: 0.8175      val loss: 0.4341
Epoch 2      train accuracy: 0.9353  train loss: 0.1697
            val accuracy: 0.8625      val loss: 0.3648
            test accuracy: 0.9025      test loss: 0.2814
```